



(a) Text chunking: one shared automaton



Total text traffic: $1 \times N$
Each byte is read once from main memory.
Parallelism comes from dividing the text.

(b) Dictionary sharding: repeated text scans



Total text traffic: $K \times N$
Shards have disjoint patterns and smaller automata.
The cost is repeated input reads plus duplicated build state.